



SF Salaries Exercise

Welcome to a quick exercise for you to practice your pandas skills! We will be using the [SF Salaries Dataset](#) from Kaggle! Just follow along and complete the tasks outlined in bold below. The tasks will get harder and harder as you go along.

Import pandas as pd.

```
In [1]: import pandas as pd
```

Read Salaries.csv as a dataframe called sal.

```
In [4]: sal = pd.read_csv('Salaries.csv')
```

Check the head of the DataFrame.

```
In [5]: sal.head()
```

Out[5]:

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	TotalPay	Tota
0	1	NATHANIEL FORD	GENERAL MANAGER-METROPOLITAN TRANSIT AUTHORITY	167411.18	0.00	400184.25	NaN	567595.43	
1	2	GARY JIMENEZ	CAPTAIN III (POLICE DEPARTMENT)	155966.02	245131.88	137811.38	NaN	538909.28	
2	3	ALBERT PARDINI	CAPTAIN III (POLICE DEPARTMENT)	212739.13	106088.18	16452.60	NaN	335279.91	
3	4	CHRISTOPHER CHONG	WIRE ROPE CABLE MAINTENANCE MECHANIC	77916.00	56120.71	198306.90	NaN	332343.61	
4	5	PATRICK GARDNER	DEPUTY CHIEF OF DEPARTMENT, (FIRE DEPARTMENT)	134401.60	9737.00	182234.59	NaN	326373.19	

Use the .info() method to find out how many entries there are.

In [6]: `sal.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148654 entries, 0 to 148653
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    148654 non-null  int64
1   EmployeeName          148654 non-null  object
2   JobTitle              148654 non-null  object
3   BasePay               148045 non-null  float64
4   OvertimePay           148650 non-null  float64
5   OtherPay              148650 non-null  float64
6   Benefits              112491 non-null  float64
7   TotalPay              148654 non-null  float64
8   TotalPayBenefits      148654 non-null  float64
9   Year                  148654 non-null  int64
10  Notes                  0 non-null       float64
11  Agency                148654 non-null  object
12  Status                 0 non-null       float64
dtypes: float64(8), int64(2), object(3)
memory usage: 14.7+ MB
```

What is the average BasePay ?

In [7]: `sal['BasePay'].mean()`

Out[7]: 66325.44884050643

What is the highest amount of OvertimePay in the dataset ?

In [8]: `sal['OvertimePay'].max()`

Out[8]: 245131.88

What is the job title of JOSEPH DRISCOLL ? Note: Use all caps, otherwise you may get an answer that doesn't match up (there is also a lowercase Joseph Driscoll).

In [9]: `sal[sal['EmployeeName']=='JOSEPH DRISCOLL']['JobTitle']`

Out[9]: 24 CAPTAIN, FIRE SUPPRESSION
Name: JobTitle, dtype: object

How much does JOSEPH DRISCOLL make (including benefits)?

In [10]: `sal[sal['EmployeeName']=='JOSEPH DRISCOLL']['TotalPayBenefits']`

Out[10]: 24 270324.91
Name: TotalPayBenefits, dtype: float64

What is the name of highest paid person (including benefits)?

In [11]: `sal[sal['TotalPayBenefits']== sal['TotalPayBenefits'].max()] #['EmployeeName']`

Out[11]:

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	TotalPay	Tota
0	1	NATHANIEL FORD	GENERAL MANAGER-METROPOLITAN TRANSIT AUTHORITY	167411.18	0.0	400184.25	NaN	567595.43	

What is the name of lowest paid person (including benefits)? Do you notice something strange about how much he or she is paid?

In [12]: `sal[sal['TotalPayBenefits'] == sal['TotalPayBenefits'].min()]`

Out[12]:

	Id	EmployeeName	JobTitle	BasePay	OvertimePay	OtherPay	Benefits	TotalPay	Tc
148653	148654	Joe Lopez	Counselor, Log Cabin Ranch	0.0	0.0	-618.13	0.0	-618.13	

What was the average (mean) BasePay of all employees per year? (2011-2014) ?

In [13]: `sal.groupby('Year').mean()['BasePay']`

Out[13]:

```
Year
2011    63595.956517
2012    65436.406857
2013    69630.030216
2014    66564.421924
Name: BasePay, dtype: float64
```

How many unique job titles are there?

In [14]: `sal['JobTitle'].nunique()`

Out[14]: 2159

What are the top 5 most common jobs?

In [15]: `sal['JobTitle'].value_counts().head(5)`

Out[15]:

```
Transit Operator    7036
Special Nurse      4389
Registered Nurse   3736
Public Svc Aide-Public Works  2518
Police Officer 3    2421
Name: JobTitle, dtype: int64
```

How many Job Titles were represented by only one person in 2013? (e.g. Job Titles with only one occurrence in 2013?)

In [16]: `sum(sal[sal['Year']==2013]['JobTitle'].value_counts() == 1)`

Out[16]: 202

How many people have the word Chief in their job title? (This is pretty tricky)

```
In [22]: def chief_string(title):  
         if 'chief' in title.lower():  
             return True  
         else:  
             return False
```

```
In [23]: sum(sal['JobTitle'].apply(lambda x: chief_string(x)))
```

Out[23]: 627

Bonus: Is there a correlation between length of the Job Title string and Salary?

```
In [20]: sal['title_len'] = sal['JobTitle'].apply(len)
```

```
In [21]: sal[['title_len', 'TotalPayBenefits']].corr()
```

Out[21]:

	title_len	TotalPayBenefits
title_len	1.000000	-0.036878
TotalPayBenefits	-0.036878	1.000000

Great Job!